

Alignment Coverage Framework

KJV 1611 / KJV 1769 Semantic Drift Study

Purpose

This document defines how the study measures and reports coverage completeness at the corpus, structural alignment, representation, and exception levels. It ensures that the paper can account for what was analyzed, what was available, what was excluded, and why.

Coverage Principles

The framework distinguishes corpus presence coverage, structural alignment coverage, and representation coverage. Corpus presence records whether works and units exist in each canon. Structural alignment records whether a shared unit_key exists in both canons. Representation coverage records whether the required named vector exists for the representation being analyzed.

Coverage Hierarchy

Coverage is computed at four levels: canon wide totals, work group totals, book level rows, and unit level exceptions. Work group reporting distinguishes Old Testament shared corpus, New Testament shared corpus, and KJV 1611 exclusive corpus where applicable.

Structural Alignment Definition

A unit is structurally aligned when the tuple (work_code, division_level_1, division_level_2) exists in both KJV 1611 and KJV 1769. Structural alignment is determined before representation eligibility. Representation coverage is therefore only meaningful after structural alignment has been established.

Representation Specific Eligibility

A structurally aligned unit is eligible for R0 drift when v_text_raw exists for both canons. It is eligible for R1 drift when v_text_norm exists for both canons. Each aligned unit receives alignment and eligibility flags so that downstream analysis can reconcile structural coverage with representation coverage.

Exception Taxonomy

Excluded or mismatched units are logged using a controlled exception taxonomy. Structural exceptions include missing sides, duplicate unit keys, invalid division keys, and non-verse units. Representation exceptions include missing named vectors, null payload text, and checksum warnings.

Study Decisions

Component	Study Decision
Structural alignment key	(work_code, division_level_1, division_level_2)

Study architecture artifact. Proprietary platform source code is not included in public release materials.

Coverage levels	Canon, work-group, book, unit-level exceptions
R0 eligibility	v_text_raw exists for both canons
R1 eligibility	v_text_norm exists for both canons
Coverage precedence	Book presence, structural alignment, R0 eligibility, R1 eligibility
Exception reporting	alignment_exceptions.jsonl

Expected Output Artifacts

- alignment_coverage_by_work.csv
- coverage_summary_table.csv
- alignment_exceptions.jsonl
- canonical_analysis_records.csv