

Corpus Scope Controls

KJV 1611 / KJV 1769 Semantic Drift Study

Purpose

This document defines the corpus scope controls used to determine which text units are eligible for direct semantic drift analysis and which units are routed into secondary descriptive outputs.

Corpus Scope Overview

The study uses two ingested and embedded corpora: KJV 1611 with 81 works and KJV 1769 with 66 works. The direct drift analysis is restricted to the shared canon intersection. KJV 1611 exclusive works are preserved as a separate corpus for future analysis.

Shared Canon Comparison Set

The shared canon comparison set is the primary analysis corpus. It consists of all work_codes that exist in both KJV 1611 and KJV 1769. This 66-work intersection is the only valid scope for direct semantic drift scoring because each unit has a possible counterpart in both editions.

KJV 1611 Exclusive Set

The KJV 1611 exclusive set consists of all work_codes present in KJV 1611 but absent from KJV 1769. This includes the apocryphal and extra-canonical material in the KJV 1611 dataset, with Baruch and the Letter of Jeremiah retained as distinct works according to the ingested dataset structure.

Exclusive-set units are not used for direct drift scoring because no aligned KJV 1769 counterpart exists. They are preserved for future nearest-neighbor mapping, thematic clustering, and semantic region analysis.

Primary Comparison Rule

A text unit is eligible for the main drift study only when its work_code belongs to the shared canon intersection, its structural unit key exists in both canons, and the required named vector exists for the representation under analysis. This produces parallel Shared Canon R0 and Shared Canon R1 comparison universes with the same structural scope.

Book Identity Policy

The ingested work_code system is the source of truth. No external canonical renaming layer is imposed after ingestion, and no books are collapsed after ingestion. Baruch and the Letter of Jeremiah remain separate works because they are represented separately in the dataset.

Coverage Accounting Requirement

Coverage reporting accounts for all works and units across both corpora. Per work reporting includes canon presence, total units, aligned units, unaligned units, and representation coverage for R0 and R1.

Study Decisions

Component	Study Decision
KJV 1611 corpus size	81 works
KJV 1769 corpus size	66 works
Primary drift corpus	66 work shared canon intersection
Exclusive corpus treatment	Inventoried and excluded from direct aligned pair drift calculations
Book identity source of truth	Ingested work code system
Representation-aware scope	R0 and R1 coverage tracked separately

Expected Output Artifacts

- alignment_coverage_by_work.csv
- coverage_summary_table.csv
- scope_registry.csv
- exclusive_dataset_inventory.csv