

Text Unit Representation Policy

KJV 1611 / KJV 1769 Semantic Drift Study

Purpose

This document defines how textual representations are constructed for embedding and drift measurement. It ensures that representation-specific outputs remain interpretable, comparable, and reproducible while supporting later analysis of orthographic and surface form effects.

Representation Types

Raw Representation (R0)

R0 uses the `text_raw` field exactly as ingested. It preserves spelling, punctuation, capitalization, historical orthography, ampersands, and character-level features. R0 is the primary reporting layer for the study because it captures the total textual form available to the embedding model.

Lightly Normalized Representation (R1)

R1 uses the `text_norm` field. The normalization profile lowercases the text, trims leading and trailing whitespace, and collapses repeated whitespace. It does not perform historical orthographic modernization or lexical mapping.

Reserved Aggressive Normalization Layer (R2)

R2 is reserved for a future study phase. It is designed to test historically aware orthographic normalization and automated lexical mapping, but it is not active in the present analysis.

Representation Policy

The study computes drift independently for each active representation. R0 and R1 use the same embedding model and vector storage architecture but remain separate representation specific statistical universes. R0 is used for primary inferential reporting. R1 is retained as a light normalization invariance and sensitivity check.

Sensitivity Interpretation

Representation sensitivity is evaluated by comparing `drift_R0` and `drift_R1` on matched units. If R0 drift is substantially higher than R1 drift, the difference is interpreted as evidence that the current normalization profile reduced surface form distance. If R0 and R1 are approximately equal, the selected normalization profile did not materially change model behavior for the corresponding unit.

Representation Provenance

Each study run records the representation type, embedding model, embedding engine, normalization profile, stable unit identifiers, and unit checksums. This metadata permits later verification that representation specific comparisons were computed from the intended text fields and embedding model.

Study Decisions

Component	Study Decision
R0 source field	text_raw
R1 source field	text_norm
R1 normalization profile	Lowercase, trim whitespace, collapse repeated whitespace
Primary reporting layer	R0
Sensitivity layer	R1
Reserved future layer	R2 aggressive historical normalization

Expected Output Artifacts

- drift_pairs.csv
- sensitivity_indices_by_scope.csv
- manifest_run.json